

612 Week 4 Homework: dplyr

Yunting

2/8/2020

Exercise 1: Pygmalion Data

```
library(Sleuth3)
library(tidyverse)

## -- Attaching packages --- tidyverse 1.3.0 --

## <U+2713> ggplot2 3.2.1      <U+2713> purrr 0.3.3
## <U+2713> tibble 2.1.3      <U+2713> dplyr 0.8.3
## <U+2713> tidyr 1.0.0       <U+2713> stringr 1.4.0
## <U+2713> readr 1.3.1      <U+2713> forcats 0.4.0

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()

data(case1302)
head(case1302)
```

```
##   Company      Treat Score
## 1      C1 Pygmalion  80.0
## 2      C1   Control  63.2
## 3      C1   Control  69.2
## 4      C2 Pygmalion  83.9
## 5      C2   Control  63.1
## 6      C2   Control  81.5
```

1. What are the units? What are the variables?

Units : platoons
Variables: Company, Treat and Score

```
library(Sleuth3)
library(tidyverse)
library(scales)
```

2. Calculate the average score within each company by treatment combination.

```
##
## Attaching package: 'scales'

## The following object is masked from 'package:purrr':
##
##     discard

## The following object is masked from 'package:readr':
##
##     col_factor
```

```
case1302 %>%
  group_by(Company,Treat) %>% #calculate the mean Score of Company and Treat
  summarize(average_Score = mean(Score, na.rm = TRUE), n= n())
```

```
## # A tibble: 20 x 4
## # Groups:   Company [10]
##   Company Treat   average_Score     n
##   <fct>   <fct>         <dbl> <int>
## 1 C1      Control          66.2     2
## 2 C1      Pygmalion          80      1
## 3 C10     Control          70.7     2
## 4 C10     Pygmalion          83.7     1
## 5 C2      Control          72.3     2
## 6 C2      Pygmalion          83.9     1
## 7 C3      Control          76.2     1
## 8 C3      Pygmalion          68.2     1
## 9 C4      Control          66.5     2
## 10 C4     Pygmalion          76.5     1
## 11 C5     Control          76.2     2
## 12 C5     Pygmalion          87.8     1
## 13 C6     Control          81.8     2
## 14 C6     Pygmalion          89.8     1
## 15 C7     Control          65.1     2
## 16 C7     Pygmalion          76.1     1
## 17 C8     Control          70.5     2
## 18 C8     Pygmalion          71.5     1
## 19 C9     Control          73.1     2
## 20 C9     Pygmalion          69.5     1
```

```
library(ggthemes)
case1302 %>%
  group_by(Company,Treat) %>%
  summarize(mean_Score = mean(Score, na.rm = TRUE), n = n()) ->


sumdf
head(sumdf)
```

3. Copy this plot:

```
## # A tibble: 6 x 4
## # Groups:   Company [3]
##   Company Treat    mean_Score    n
##   <fct>   <fct>      <dbl> <int>
## 1 C1      Control      66.2     2
## 2 C1      Pygmalion      80      1
## 3 C10     Control      70.7     2
## 4 C10     Pygmalion      83.7     1
## 5 C2      Control      72.3     2
## 6 C2      Pygmalion      83.9     1
```

```
sumdf %>%
  ggplot(aes(x = Company, y = mean_Score)) +
  theme_bw() +
  geom_point(aes(size = n, color = Treat)) +
  scale_color_colorblind(name = "Treatment") +
  scale_size(name = "Number of \nObersavtions", range = c(2,4), breaks = c(1,2)) +
  guides("Treatment", "Number of Obersavtions") +
  ylab("Mean Score") +
  guides(color = guide_legend(order=1), size = guide_legend(order=2))
```

```
## Warning: Duplicated aesthetics after name standardisation: NA
```




```
library(ggthemes)
case1302 %>%
  group_by(Company,Treat) %>%
  summarize(mean_Score = mean(Score, na.rm = TRUE), n = n()) %>%

  ggplot(aes(x = Company, y = mean_Score, size = n, color = Treat)) +
  theme_bw() +
  geom_point() +
  scale_color_colorblind(name = "Treatment") +

  scale_size(name = "Number of \nObersavtions", range = c(2,4), breaks = c(1,2)) +
  guides("Treatment", "Number of Obersavtions") +
  ylab("Mean Score") +
  guides(color = guide_legend(order=1),size = guide_legend(order=2))
```

Warning: Duplicated aesthetics after name standardisation: NA



4. Does it make sense to add a loess smoother to the above plot? Why or why not? If so, add one.

I suggest we cannot add a `geom_smooth` line. Although we added the function, the spot of this chart is too separate. In this case, it is not make sense to add a loess smoother to the above plot.

real ans: because a loess smoother would only make sense if the explanatory variable was also quantitative, but it is categorical

Exercise 2: Midwest Data

- For this exercise, we'll use the midwest data from ggplot2.

```
data(midwest)
head(midwest)
```

1. Load these data into R.

```
## # A tibble: 6 x 28
##   PID county state area poptotal popdensity popwhite popblack popamerindian
##   <int> <chr> <chr> <dbl> <int> <dbl> <int> <int> <int>
## 1  561 ADAMS IL 0.052 66090 1271. 63917 1702 98
## 2  562 ALEXA... IL 0.014 10626 759 7054 3496 19
## 3  563 BOND IL 0.022 14991 681. 14477 429 35
## 4  564 BOONE IL 0.017 30806 1812. 29344 127 46
## 5  565 BROWN IL 0.018 5836 324. 5264 547 14
## 6  566 BUREAU IL 0.05 35688 714. 35157 50 65
## # ... with 19 more variables: popasian <int>, popother <int>, percwhite <dbl>,
## # percblack <dbl>, percamerindian <dbl>, percasian <dbl>, percother <dbl>,
## # popadults <int>, perchsd <dbl>, percollege <dbl>, percprof <dbl>,
## # poppovertyknown <int>, percpovertyknown <dbl>, percbelowpoverty <dbl>,
## # percchildbelowpovert <dbl>, percadultpoverty <dbl>,
## # percelderlypoverty <dbl>, inmetro <int>, category <chr>
```


2. What are the observational units?

Ans: counties


```
midwest %>%
  group_by(state) %>%
  mutate(sum_poptotal = sum(poptotal, na.rm = TRUE))%>%

ggplot(aes(x = state, y = sum_poptotal)) +
  geom_col()
```

3. Calculate the total population in each state. Display the data in an appropriate plot




```
midwest %>%  
  group_by(state) %>%  
  mutate(sum_poptotal = sum(poptotal, na.rm = TRUE)) %>%  
ggplot(aes(x = state, y = sum_poptotal))+  
  theme_economist_white()+  
  geom_col()
```



```
midwest %>%
  filter(poptotal > quantile(poptotal, probs = 0.1)) %>%
  ggplot( mapping = aes(x = popdensity, y = percollege, color = state)) +
  geom_point(size = 1) +
  geom_smooth(method = "lm", se = FALSE, linetype = "dotted", color = "black")+
  theme_bw() +
  scale_x_log10() +
  scale_y_log10()
```


4. Make a scatterplot of population density against the percent college educated on a scale that is appropriately linear. You should exclude counties in the bottom tenth percentile of total population. Color code the counties by state. Add the overall OLS line (the one that takes all points into account). Make the OLS line black and dotted.



```
midwest %>%
  mutate(percadults = popadults / poptotal) %>%
  arrange(desc(percadults)) %>%

  ggplot(mapping = aes(x = state, y = percadults, fill = state)) +
  theme_bw() +
  geom_boxplot() +
  ylab("Percentage of Adults") +
  xlab("States of midwest")
```

5. Make an appropriate plot to explore the association between the state and the percentage adults. Make sure the axis labels are nice and use the black-and-white theme.



```
unique(midwest$state)
```

6. Use an R function to determine the possible values of state.

```
## [1] "IL" "IN" "MI" "OH" "WI"
```

```
midwest %>%
  mutate(state = recode(state, IL = "Illinois", IN = "Indiana", MI = "Michigan", OH = "Ohio", WI = "Wisconsin"))
tail()
```

7. In the state variable, replace or recode the abbreviations with the full state name, e.g., IL with Illinois, IN with Indiana, etc..

```
## # A tibble: 6 x 28
##   PID county state area poptotal popdensity popwhite popblack popamerindian
##   <int> <chr> <chr> <dbl> <int> <dbl> <int> <int> <int>
## 1 3047 WASHI... Wisc... 0.025 95328 3813. 94465 125 208
```

```
## 2 3048 WAUKE... Wisc... 0.034 304715 8962. 298313 1096 672
## 3 3049 WAUPA... Wisc... 0.045 46104 1025. 45695 22 125
## 4 3050 WAUSH... Wisc... 0.037 19385 524. 19094 29 70
## 5 3051 WINNE... Wisc... 0.035 140320 4009. 136822 697 685
## 6 3052 WOOD Wisc... 0.048 73605 1533. 72157 90 481
## # ... with 19 more variables: popasian <int>, popother <int>, percwhite <dbl>,
## # percblack <dbl>, percamerindan <dbl>, percasian <dbl>, percother <dbl>,
## # popadults <int>, perchsd <dbl>, percollege <dbl>, percprof <dbl>,
## # poppovertyknown <int>, percpovertyknown <dbl>, percbelowpoverty <dbl>,
## # percchildbelowpovert <dbl>, percadultpoverty <dbl>,
## # percelderlypoverty <dbl>, inmetro <int>, category <chr>
```

Exercise 3: Coding Problems

```
lsumfun <- function(x){
  stopifnot(is.numeric(x))
  if (length(x) %% 2 == 0){
    return(sum(x[x], na.rm = TRUE))
  } else {
    return(sum(x[x %% 2 == 0], na.rm = TRUE))
  }
}

lsumfun(c(1, 2, 3, NA))
```

1. Create a function that takes a vector of numerics as input. It checks the length of the vector. If the length is even, it returns the sum of the vector. If the length is odd, it returns the sum of all of the even numbers of the vector. For example, the following are some outputs of one implementation, called `lsumfun()`.

```
## [1] 6
```

```
lsumfun(c(1, 2, 3))
```

```
## [1] 2
```

```
lsumfun(c(2, 3, 4, 5, 6))
```

```
## [1] 12
```

```
lsumfun(c(2, 3, 4, 5, NA))
```

```
## [1] 6
```

```
lsumfun(c(2, 3, 4, NA, 6))
```

```
## [1] 12
```

2. We add a Leap Day on February 29, almost every four years. In the Gregorian calendar three criteria must be taken into account to identify leap years:

- The year can be evenly divided by 4, is a leap year, unless:
- The year can be evenly divided by 100, it is NOT a leap year, unless:
- The year is also evenly divisible by 400. Then it is a leap year.

This means that in the Gregorian calendar, the years 2000 and 2400 are leap years, while 1800, 1900, 2100, 2200, 2300 and 2500 are NOT leap years

Write a function that takes the year as input and returns TRUE if it is a leap year and FALSE if it is not a leap year. Evaluate your function at 2, 12, 200, 800.

```
Gregorian_calendar <- function(y){  
  if (y %% 400 == 0){return(" TRUE")}  
  else if (y %% 100 ==0){return(" FALSE")}  
  else if (y %% 4 ==0){return(" TRUE")}  
  
  else {return("FALSE")}  
  
}  
Gregorian_calendar(2)
```

```
## [1] "FALSE"
```

```
Gregorian_calendar(12)
```

```
## [1] " TRUE"
```

```
Gregorian_calendar(200)
```

```
## [1] " FALSE"
```

```
Gregorian_calendar(800)
```

```
## [1] " TRUE"
```

3. Create a function that takes two vectors of numerics as input, a and b.

- Your vector should throw an error if either a or b contain repeated elements.
- If a number is in a and not b you add 1 to your score.
- If a number is in b and not a you subtract 1 from your score.
- If a number is in both a and b you do nothing to your score.
- The function returns the final score.
- Hint: `help("%in%")`

```
score2 <- function(a,b){  
  stopifnot(is.numeric(a))  
  stopifnot(is.numeric(b))
```

```

if(length(a) != length(unique(a))){
  stop("a has some repeated elements ")
}else if(length(b) != length(unique(b))){
  stop("b has some repeated elements")
}

return(sum(!(a %in% b)) - sum(!(b %in% a)))
}

score2(c(1, 2, 3), c(3, 4))

```

```
## [1] 1
```

```
score2(c(1, 2, 3), c(1, 2, 3, 4))
```

```
## [1] -1
```

```
score2(c(1, 2, 3), c(3, 4, 4))
```

```
## Error in score2(c(1, 2, 3), c(3, 4, 4)): b has some repeated elements
```